

Progress Report

Which tasks have been completed?

1. We created a StackOverflow crawler and successfully fetched over 1,000 training examples of StackOverflow question-answer pairs that link to Python documentation.
2. We indexed Elasticsearch on all of the subpages of <https://docs.python.org/3/index.html> and wrote the evaluator logic that computes the reciprocal rank of the ground truth answer for a StackOverflow question as retrieved by Elasticsearch.
3. We ran GENRE end-to-end on a single training example.

Which tasks are pending?

1. Run our StackOverflow crawler for an extended period of time to fetch all training data. We estimate that we will have around 30,000 training examples.
2. Split fetched question-answer pairs into training and testing data sets.
3. Train GENRE on the training data set. Create evaluation logic for GENRE.
4. Run comparisons of GENRE and Elasticsearch performance.

Are you facing any challenges?

No, but we are cognizant of the risks of not having enough training data and/or not being able to train GENRE properly. These are the two areas we are prioritizing and are aiming to de-risk as soon as possible.